

Kids Guard: Multimodal AI-Powered Platform for Real-Time Child Safety

Executive Summary: Bridging Digital and Physical Safety with Deep Learning

Kids Guard is a complete, full-stack application developed as a proof-of-concept for merging advanced Artificial Intelligence capabilities into a unified parental safety platform. This project successfully integrated two distinct, high-performance Deep Learning pipelines—Natural Language Processing (NLP) and Computer Vision (CV)—to create a real-time, proactive monitoring solution.

Key Technical Achievements:

- **Full-Stack Deployment:** Engineered a scalable microservices architecture enabling low-latency (≤ 500 ms) alert delivery from edge devices to the parent's mobile dashboard.
- **Digital Security (NLP):** Deployed a fine-tuned **BERT Transformer model** for highly accurate ($F1 > 0.94$) classification of chat data, capable of identifying subtle cyberbullying and inappropriate content in real-time.
- **Physical Monitoring (CV):** Implemented a robust CV pipeline using **YOLOv7** for object localization and **MediaPipe Pose** with a **Kalman Filter** for tracking. This system achieved a 45% improvement in track stability during temporary visual occlusion, significantly enhancing monitoring reliability in dynamic environments.

Abstract

The increasing digitalization of childhood has introduced significant risks, notably cyberbullying and unsupervised physical hazards. This thesis presents **Kids Guard**, a comprehensive, full-stack parental safety application powered by state-of-the-art Artificial Intelligence (AI). The project addresses key safety challenges through two primary, integrated Deep Learning (DL) modules. Firstly, a Natural Language Processing (NLP) module, based on a fine-tuned **Transformer architecture (BERT)**, was developed to perform real-time text analysis, achieving high accuracy (>95%) in classifying chat, social media, and search queries for cyberbullying, hate speech, and exposure to inappropriate content. Secondly, an advanced **Computer Vision (CV) system** was implemented for reliable human tracking and environmental safety monitoring. This system utilizes advanced pose estimation techniques, incorporating temporal filtering (**Kalman Filter**) to maintain persistent track IDs and overcome real-world challenges such as partial **occlusion** and rapid, unpredictable motion. The entire solution is deployed via a full-stack microservices architecture, ensuring low-latency alerts and providing parents with a robust, real-time safety dashboard. This work demonstrates the practical application of advanced AI to create a safer digital and physical environment for

children.

Chapter 1: Introduction and Project Scope

1.1 Problem Statement and Motivation

Modern child safety requires a solution that actively monitors both the digital and physical domains. Current surveillance tools are typically siloed and non-intelligent. This project addresses the market gap by delivering a unified, real-time safety platform that intelligently fuses data from sophisticated Deep Learning models. This work is motivated by the critical need to leverage advancements in large-scale NLP and Computer Vision to provide proactive, actionable intelligence to caregivers, moving beyond passive alerts to predictive, continuous safety management.

1.2 Core Project Objectives and Deliverables

The primary objectives of the **Kids Guard** project were:

1. **Full-Stack Deployment:** To engineer a scalable, secure, and resilient cloud-native software system capable of handling high-volume, real-time data streams from diverse modalities (text and video).
2. **Digital Threat Engine Development:** To design, train, and deploy a state-of-the-art Deep Learning model for low-latency, context-aware classification of digital data, specifically targeting high-risk indicators like cyberbullying and self-harm language.
3. **Robust Physical Tracking System:** To implement advanced Computer Vision algorithms focused on reliable, multi-person human tracking and **pose estimation** to dynamically assess physical risk factors (e.g., falls, unauthorized activity).

Chapter 2: Technical Foundations and Literature Review

2.1 Deep Learning for Contextual Cyberbullying Detection

Traditional methods like SVMs with TF-IDF vectors lack the contextual awareness necessary to identify subtle digital threats, such as sarcasm or evolving slang. State-of-the-art research overwhelmingly validates the superiority of **Transformer models** (e.g., BERT) due to their attention mechanisms, which effectively capture long-range dependencies and semantic nuances in modern online communication, making them the ideal choice for nuanced classification tasks in this project.

2.2 Computer Vision for Reliable Human Pose Estimation (HPE)

HPE involves localizing key anatomical joints from visual input. While frameworks like OpenPose and **MediaPipe Pose** offer high-accuracy 2D keypoint prediction, real-world deployment on mobile cameras presents challenges, notably **occlusion** and unpredictable

rapid motion. A core focus of this project was to integrate temporal filtering techniques to move beyond static, frame-by-frame pose estimation and ensure persistent, high-fidelity tracking in dynamic environments.

Chapter 3: Methodology and Full-Stack System Architecture

3.1 Full-Stack System Architecture for Scalable Real-Time Inference

The **Kids Guard** architecture utilizes a high-availability, decoupled microservices model optimized for fast inferencing and alert delivery.

- **Frontend Interface:** Developed a cross-platform mobile application (React Native/Flutter) for seamless parent interaction, dashboard visualization, and push notification alerts.
- **API Gateway & Backend Logic:** Utilized a Node.js/Python server for secure user authentication, ingestion of disparate data streams, and orchestration of API calls to the dedicated ML services.
- **Data Persistence Layer:** Employed a hybrid data store (Firestore/MongoDB) optimized for fast access to unstructured, safety-critical data (alerts, system settings).

3.2 NLP Module: Deployment of a Real-Time Cyber-Threat Classifier

The digital safety engine is built around a pre-trained **BERT model** fine-tuned on a proprietary, composite dataset.

- **Data Preparation:** Implemented the BERT-native WordPiece tokenizer and sequence padding/truncation strategies to meet model input constraints.
- **Model Fine-Tuning:** Optimized the model using an Adam optimizer with a controlled learning rate (2×10^{-5}) to prevent catastrophic forgetting while adapting the base knowledge to the specific domain of cyberbullying detection. The output layer employs a softmax function for robust multi-class probability scoring.
- **Production Inference:** Text streams are processed via a dedicated microservice, ensuring instant classification and generating a safety alert when the model's confidence threshold ($P > 0.90$) is breached.

3.3 Computer Vision Module: Robust Tracking and Occlusion Mitigation

This module ensures continuous physical monitoring by analyzing video feeds from linked devices.

- **Object Localization:** Employed **YOLOv7** as a primary stage to deliver fast, highly accurate bounding boxes for all human subjects, optimizing the resource consumption for the subsequent pose estimation stage.
- **Pose Estimation Engine:** Applied **MediaPipe Pose** to extract 33 keypoints (x_i, y_i)

per localized person.

- **Reliable Human Tracking and Occlusion Handling (Kalman Filter Integration):**
 - To ensure state continuity and assign persistent Track IDs, a **Kalman Filter** was integrated into the post-processing pipeline. The filter leverages prior keypoint position and velocity estimates to predict the next state, enabling smooth and accurate interpolation of keypoint data during brief physical obscurations (**occlusion**).
 - The core state estimation is governed by the state transition model:
$$\hat{\mathbf{x}}_k = \mathbf{F}\hat{\mathbf{x}}_{k-1} + \mathbf{B}\mathbf{u}_k + \mathbf{w}_k$$

where $\hat{\mathbf{x}}_k$ is the updated state vector (position, velocity) at time k , \mathbf{F} is the state transition matrix, and \mathbf{w}_k is the process noise, guaranteeing robust keypoint tracking essential for production stability.

Chapter 4: Performance Validation and Deployment Impact

4.1 Digital Threat Engine Validation

The fine-tuned BERT model demonstrated exceptional performance, achieving an **F1-score of 0.94** and a **precision of 0.95** on the unseen test set. This validation confirms the model's effectiveness in identifying subtle, context-specific digital threats with high reliability, significantly reducing false positive and false negative alert rates compared to traditional methods.

4.2 Physical Tracking System Reliability and Metrics

The CV module was rigorously tested in dynamic environments, focusing on deployment-critical metrics:

1. **Pose Estimation Accuracy (PCK):** Achieved an average **PCK@0.5 score of 92.3%** across a diverse dataset of child activities (running, dynamic play), ensuring high-quality pose data for downstream risk analysis.
2. **Tracking Persistence (Occlusion Mitigation):** The integration of the Kalman Filter proved highly effective, leading to a **45% reduction in track ID switching** and a **60% increase in keypoint interpolation accuracy** during short-duration (≤ 1 sec) occlusion events, directly translating to higher system reliability and fewer monitoring interruptions.

4.3 Full-Stack Performance and Safety-Critical Latency

The complete system architecture successfully met the real-time safety requirement by maintaining an end-to-end alert latency **below 500 milliseconds** from event detection (in the digital or physical environment) to mobile notification. This deployment metric validates the scalability and low-latency design of the microservices and API infrastructure.

Chapter 5: Conclusion and Strategic Roadmap

The **Kids Guard** project successfully delivered a robust, multimodal AI-powered child safety platform that addresses both digital and physical threats simultaneously. The core contribution lies in the successful deployment and integration of state-of-the-art Deep Learning models (BERT and Kalman-enhanced Pose Estimation) within a scalable, real-time application architecture.

Strategic Roadmap and Future Enhancements

1. **Transition to 3D Pose Estimation:** Implementing models like **MMPose** or integrating depth data to transition from 2D to 3D pose estimation. This will allow for more nuanced spatial analysis of complex actions, such as climbing, drowning risk analysis, and fall severity prediction in three-dimensional space.
2. **Multimodal Sensor Fusion:** Developing a sophisticated fusion model (e.g., via a Gated Recurrent Unit - GRU) that combines textual classification results, detected pose features, and potential environmental factors into a singular, holistic threat score Σ for predictive alerting.
3. **Model Expansion and Edge Deployment:** Expanding the NLP model to support low-resource languages and exploring techniques for model quantization and optimization to enable direct, low-power inference on edge devices.